



Network Intrusion Detection - Apache Spark


Team 2 - Shubham Chawla, Chintan Vajani, Archit Sinha

Project Summary

A system's security and privacy are jeopardised when an intrusion occurs. The Intrusion Detection System (IDS) is essential for network security since it can identify various kinds of attacks. Therefore, utilising machine learning algorithms, we will propose an intrusion detection system here. The suggested system will be tested using the KDD99 Dataset. Our project includes implementation and performance evaluation of a Big Data project whose goal is to use Apache Spark and MLlib in order to perform network traffic analysis on KDD99 dataset. Further the project was deployed on Flask and hosted on Heroku.

Dataset link: <http://kdd.ics.uci.edu/databases/kddcup99/kddcup99.html>
(<http://kdd.ics.uci.edu/databases/kddcup99/kddcup99.html>)

Data Mining Process



Ref: https://en.wikipedia.org/wiki/Cross-industry_standard_process_for_data_mining
https://en.wikipedia.org/wiki/Cross-industry_standard_process_for_data_mining

The process of generating and implementing data-driven solutions to business problems is iterative, and the model highlights this by highlighting a number of distinct stages that are frequently revisited during this process.

- Business understanding
- Data understanding
- Data exploration and preparation
- Model Building
- Deployment

Business Understanding

Numerous firms are now at danger as a result of the growing cyber security risks. Any organisation could be the target of attackers who can cause significant harm to it.

For these reasons, it's essential to offer a system that aids in the detection of network intrusions. In order to facilitate the deployment and to obtain findings as quickly as possible, this intrusion detection system should be non-invasive to other systems.

Data Understanding

The dataset includes samples of network connections that are either normal or under attack; in this instance, the

attack type is mentioned. A connection is a sequence of TCP packets starting and ending at some well defined times, between which data flows to and from a source IP address to a target IP address under some well defined protocol. Each connection is labeled as either normal, or as an attack, with exactly one specific attack type. Each connection record consists of about 100 bytes.

It consists of several attack methods like Denial of Service (DoS), port scanning, and buffer overflow along with normal samples that have all been combined into a single comma-separated values (CSV) file with 4898431 samples and 42 characteristics. The resulting features are categorised into the following groups:

Attacks fall into four main categories:

- **DOS:** denial-of-service, e.g. syn flood
- **R2L:** unauthorized access from a remote machine, e.g. guessing password
- **U2R:** unauthorized access to local superuser (root) privileges, e.g., various ``buffer overflow" attacks
- **probing:** surveillance and other probing, e.g., port scanning.

Features essential for distniguishing *normal* connections from *attacks*

Basic Features of each network connection vector

No	Feature name	description	type
1	duration	length (number of seconds) of the connection	continuous
2	protocol_type	type of the protocol, e.g. tcp, udp, etc.	discrete
3	service	network service on the destination, e.g., http, telnet, etc.	discrete
4	flag	normal or error status of the connection. The possible status are this: SF, S0, S1, S2, S3, OTH, REJ, RSTO, RSTOS0, SH, RSTRH, SHR	discrete
5	src_bytes	number of data bytes from source to destination	continuous
6	dst_bytes	number of data bytes from destination to source	continuous
7	land	1 if connection is from/to the same host/port; 0 otherwise	discrete
8	wrong_fragment	sum of bad checksum packets in a connection	continuous
9	urgent	number of urgent packets. Urgent packets are packets with the urgent bit activated	continuous

Content Related Features of each network connection vector

No.	Feature name	description	type
10	hot	sum of hot actions in a connection such as: entering a system direc- tory, creating programs and executing programs	continuous
11	num_failed_logins	number of failed login attempts	continuous
12	logged_in	1 if successfully logged in; 0 otherwise	discrete
13	num_compromised	number of "compromised" conditions	continuous
14	root_shell	1 if root shell is obtained; 0 otherwise	discrete
15	su_attempted	1 if "su root" command attempted; 0 otherwise	discrete
16	num_root	number of "root" accesses	continuous

No.	Feature name	description	type
17	num_file_creations	number of file creation operations	continuous
18	num_shells	number of shell prompts	continuous
19	num_access_files	number of operations on access control files	continuous
20	num_outbound_cmds	number of outbound commands in an ftp session	continuous
21	is_hot_login	1 if the login belongs to the "hot" list; 0 otherwise	discrete
22	is_guest_login	1 if the login is a "guest" login; 0 otherwise	discrete

Time Related Traffic Features of each network connection vector

Col	Feature name	description	type
23	count	sum of connections to the same destination IP address	continuous
24	srv_count	sum of connections to the same destination port number	continuous
25	serror_rate	the percentage of connections that have activated the flag (4) s0, s1, s2 or s3, among the connections aggregated in count (23)	continuous
26	srv_serror_rate	the percentage of connections that have activated the flag (4) s0, s1, s2 or s3, among the connections aggregated in srv_count (24)	continuous
27	rerror_rate	the percentage of connections that have activated the flag (4) REJ, among the connections aggregated in count (23)	continuous
29	same_srv_rate	the percentage of connections that have activated the flag (4) REJ, among the connections aggregated in srv_count (24)	continuous
28	srv_rerror_rate	the percentage of connections that were to the same service, among the connections aggregated in count (23)	continuous
30	diff_srv_rate	the percentage of connections that were to different services, among the connections aggregated in count (23)	continuous
31	srv_diff_host_rate	the percentage of connections that were to different destination ma- chines among the connections aggregated in srv_count (24)	continuous

Host Based Traffic Features of each network connection vector

Col	Feature name	description	type
32	dst_host_count	sum of connections to the same destination IP address	continuous
33	dst_host_srv_count	sum of connections to the same destination port number	continuous
34	dst_host_same_srv_rate	the percentage of connections that were to the same service, among the connections aggregated in dst_host_count (32)	continuous
35	dst_host_diff_srv_rate	the percentage of connections that were to different services, among the connections aggregated in dst_host_count (32)	continuous
36	dst_host_same_src_port_rate	the percentage of connections that were to the same source port, among the connections aggregated in dst_host_srv_count (33)	continuous
37	dst_host_srv_diff_host_rate	the percentage of connections that were to different destination ma- chines, among the connections aggregated in dst_host_srv_count (33)	continuous
38	dst_host_serror_rate	the percentage of connections that have activated the flag (4) s0, s1, s2 or s3, among the connections aggregated in dst_host_count (32)	continuous
39	dst_host_srv_serror_rate	the percent of connections that have activated the flag (4) s0, s1, s2 or s3, among the connections aggregated in dst_host_srv_count (33)	continuous
40	dst_host_rerror_rate	the percentage of connections that have activated the flag (4) REJ, among the connections aggregated in dst_host_count (32)	continuous

Col	Feature name	description	type
41	dst_host_srv_error_rate	the percentage of connections that have activated the flag (4) REJ, among the connections aggregated in dst_host_srv_count (33)	continuous

The 42nd attribute is our target variable where we will be using Multinomial Classification to predict whether our class is Normal or DOS or PROBE or R2L or U2R

Initializing Apache Spark and Importing Necessary Libraries

```
In [1]: import pyspark
        from pyspark.context import SparkContext
        import warnings
        warnings.filterwarnings("ignore")
        import time
        from pyspark.sql import SparkSession

        # To perform feature engineering
        from pyspark.ml.feature import StringIndexer, VectorAssembler
        from pyspark.ml import Pipeline

        # To classify label class
        from pyspark.ml.classification import LogisticRegression, DecisionTreeClassifier, \
            NaiveBayes, RandomForestClassifier

        # Importing MulticlassClassificationEvaluator to evaluate the performance of the classifiers (models),
        from pyspark.ml.evaluation import MulticlassClassificationEvaluator

        from pyspark.sql.functions import regexp_replace,lit
        # Spark Session
        from pyspark.sql import SparkSession
        spark = SparkSession.builder.getOrCreate()
        spark.conf.set('spark.sql.shuffle.partitions', 6)
        spark.conf.set('num-executors', 16)
        spark

SparkSession - in-memory

SparkContext

Spark UI (http://ShubhamPC:4040)

Version
v3.3.0
Master
local[*]
AppName
pyspark-shell
```

Data Exploration & Preparation

```

In [2]: import seaborn as sns
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline

In [3]: # Defining attribute names as they were not present in the data file
columns = ["duration", "protocol_type", "service", "flag", "src_bytes", "dst_bytes",
           "land", "wrong_fragment", "urgent", "hot", "num_failed_logins", "logged_in",
           "num_compromised", "root_shell", "su_attempted", "num_root", "num_file_creations",
           "num_shells", "num_access_files", "num_outbound_cmds", "is_host_login",
           "is_guest_login", "count", "srv_count", "serror_rate", "srv_serror_rate", "rerror_rate",
           "srv_rerror_rate", "same_srv_rate", "diff_srv_rate", "srv_diff_host_rate", "dst_host_count",
           "dst_host_srv_count", "dst_host_same_srv_rate", "dst_host_diff_srv_rate",
           "dst_host_same_src_port_rate", "dst_host_srv_diff_host_rate", "dst_host_serror_rate",
           "dst_host_srv_serror_rate", "dst_host_rerror_rate", "dst_host_srv_rerror_rate", "label"]

# Reading csv file into spark dataframe
sparkdf = spark.read.csv("kddcup.data.corrected", inferSchema=True, header=False)

# Adding the list of features to our dataset
sparkdf = sparkdf.toDF(*columns)
sparkdf = sparkdf.withColumn("label", regexp_replace("label", "\.", ""))

print("Dataset sizes: {row} rows, {cols} columns".format(row=sparkdf.count(), cols=len(sparkdf.columns)))

```

Dataset sizes: 4898431 rows, 42 columns

<https://stackoverflow.com/questions/39067505/pyspark-display-a-spark-data-frame-in-a-table-format>
<https://stackoverflow.com/questions/39067505/pyspark-display-a-spark-data-frame-in-a-table-format>
[\(https://sparkbyexamples.com/pyspark/pyspark-find-count-of-null-none-nan-values/#:~:text=In%20PySpark%20DataFrame%20you%20can,count\(\)%20and%20when\(\)\)](https://sparkbyexamples.com/pyspark/pyspark-find-count-of-null-none-nan-values/#:~:text=In%20PySpark%20DataFrame%20you%20can,count()%20and%20when())
[\(https://sparkbyexamples.com/pyspark/pyspark-find-count-of-null-none-nan-values/#:~:text=In%20PySpark%20DataFrame%20you%20can,count\(\)%20and%20when\(\)\)](https://sparkbyexamples.com/pyspark/pyspark-find-count-of-null-none-nan-values/#:~:text=In%20PySpark%20DataFrame%20you%20can,count()%20and%20when()).

```
In [4]: # To convert large dataframes to pandas
spark.conf.set("spark.sql.execution.arrow.pyspark.enabled", "true")

# Checking nulls in all the columns
from pyspark.sql.functions import col,isnan, when, count
sparkdf.select([count(when(isnan(c) | col(c).isNull(), c)).alias(c) for c in sparkdf.columns]
               ).toPandas()
```

	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	wrong_fragment	urgent	hot	...	dst_host_srv_count	c
0	0	0	0	0	0	0	0	0	0	0	...	0	0

1 rows x 42 columns

As we can see, there are no null values present in the dataset. So there is no need of cleaning with respect to null or missing values.

Using Pandas API on Spark to explore data further

```
In [5]: import pyspark.pandas as ps
psdf = sparkdf.to_pandas_on_spark()
```

WARNING:root:'PYARROW_IGNORE_TIMEZONE' environment variable was not set. It is required to set this environment variable to '1' in both driver and executor sides if you use pyarrow>=2.0.0. pandas-on-Spark will set it for you but it does not work if there is a Spark context already launched.

```
In [6]: type(psdf)
```

pyspark.pandas.frame.DataFrame

```
In [7]: psdf.head()
```

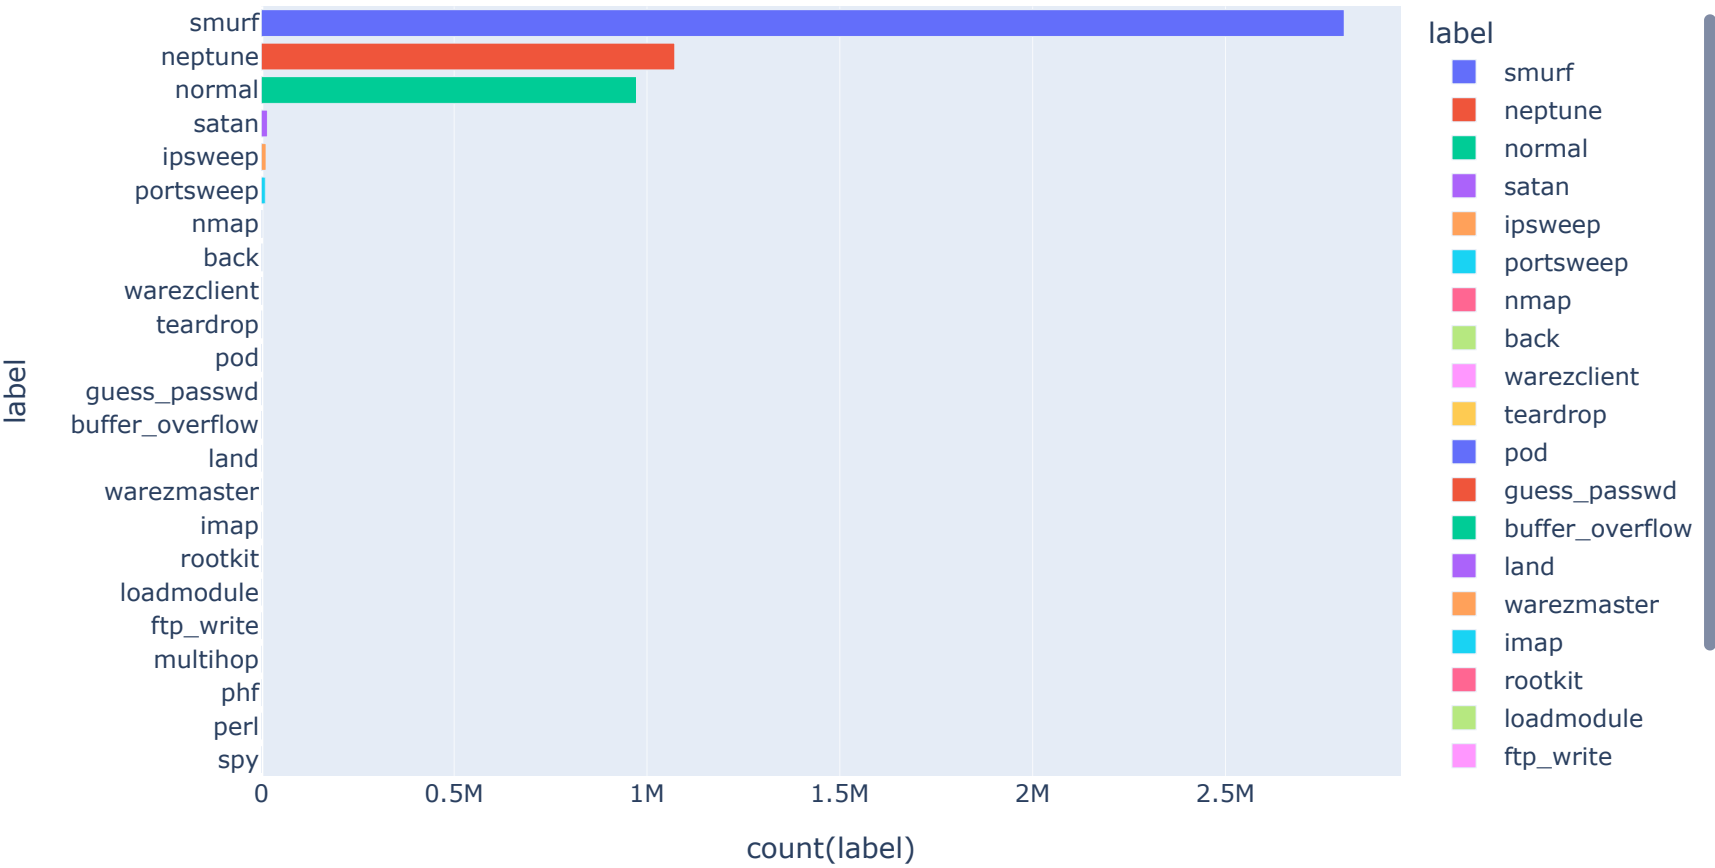
	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	wrong_fragment	urgent	hot	num_failed_logins	logged_in
0	0	tcp	http	SF	215	45076	0	0	0	0	0	1
1	0	tcp	http	SF	162	4528	0	0	0	0	0	1
2	0	tcp	http	SF	236	1228	0	0	0	0	0	1
3	0	tcp	http	SF	233	2032	0	0	0	0	0	1
4	0	tcp	http	SF	239	486	0	0	0	0	0	1


```
In [8]: # Distribution of label - attack type class
from pyspark.sql import functions as F
df1 = sparkdf.groupby('label').agg(F.count('label')).orderBy('count(label)', ascending= False)
df1.show(23)
```

label	count(label)
smurf	2807886
neptune	1072017
normal	972781
satan	15892
ipsweep	12481
portsweep	10413
nmap	2316
back	2203
warezclient	1020
teardrop	979
pod	264
guess_passwd	53
buffer_overflow	30
land	21
warezmaster	20
imap	12
rootkit	10
loadmodule	9
ftp_write	8
multihop	7

Plotting Horizontal Bar Chart for all label categories

```
In [9]: psdf = df1.to_pandas_on_spark()
psdf.plot.bar(y = 'label', x = 'count(label)', color = 'label').update_yaxes(categoryorder="total ascending")
```



There are way too many categories in the label to predict. It is better to encode them according to 5 main categories - Normal, DOS, Probe, R2L and U2R

Using Pyspark sql to encode 23 different label categories into 5 categories

- 0 - Normal
- 1 - DOS
- 2 - Probe
- 3 - R2L

• 4 - U2R

```
In [10]: # Using pyspark sql to encode attack categories into 5 different categories
#0- normal, 1- DOS, 2- Probe, 3 - R2L(remote to user attack), 4 - U2R(User to Root attack)
# Creating a temp view
sparkdf.createOrReplaceTempView("network")

# adding label_num column
df_sql = spark.sql("""
SELECT *,
CASE
    WHEN label = 'normal' THEN 0
    WHEN label = 'back' OR label = 'smurf' OR label = 'pod' OR label = 'land' OR label = 'neptune' OR label
    WHEN label = 'satan' OR label = 'ipsweep' OR label = 'nmap' OR label = 'portsweep' THEN 2
    WHEN label = 'guess_passwd' OR label = 'ftpwrite' OR label = 'imap' OR label = 'phf' OR label = 'multihop'
    ELSE 4
END AS label_num,
CASE
    WHEN label = 'normal' THEN 'normal'
    WHEN label = 'back' OR label = 'smurf' OR label = 'pod' OR label = 'land' OR label = 'neptune' OR label
    WHEN label = 'satan' OR label = 'ipsweep' OR label = 'nmap' OR label = 'portsweep' THEN 'Probe'
    WHEN label = 'guess_passwd' OR label = 'ftpwrite' OR label = 'imap' OR label = 'phf' OR label = 'multihop'
    ELSE 'U2L'
END AS label_cat
FROM network;
""")
df_sql.select(['label_num', 'label_cat']).distinct().orderBy('label_num', ascending=True).show()
```

+-----+-----+	
label_num	label_cat
+-----+-----+	
0	normal
1	DOS
2	Probe
3	R2L
4	U2L
+-----+-----+	


```
In [11]: # Checking count of new label categories
df_sql.createOrReplaceTempView("label")

df_sql1 = spark.sql("""
SELECT label_cat, count(label_cat)
FROM label
group by label_cat
order by count(label_cat)
;
""")
#df_sql.select(['label_cat']).distinct().orderBy('label_num',ascending=True).show()
df_sql1.show()
```

label_cat	count(label_cat)
R2L	98
U2L	1080
Probe	41102
normal	972781
DOS	3883370

Plotting Bar Chart for encoded label categories

```
In [12]: psdf1 = df_sql1.to_pandas_on_spark()
psdf1.plot.bar(y = 'count(label_cat)', x = 'label_cat', color = 'label_cat', text = 'count(label_cat)').update(
```



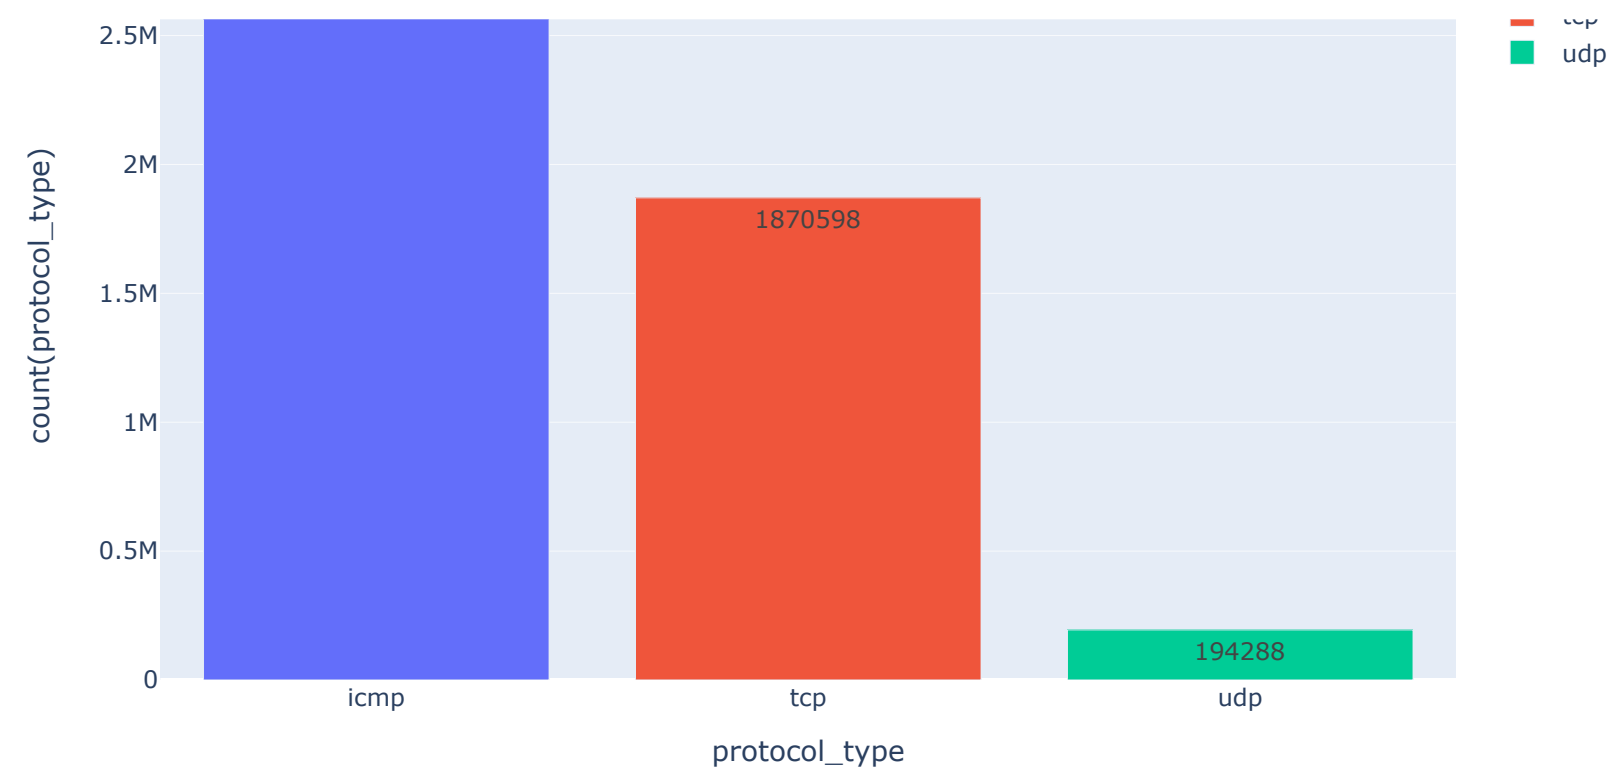
As we can see label of category 1-DOS is present in most of the data followed by normal category.

Plotting Bar Graph for different protocol types

```
In [13]: ### Plotting Bar chart for protocol types

df3 = df_sql.groupby('protocol_type').agg(F.count('protocol_type')).orderBy('protocol_type',
                                                                              ascending= True)


psdf2 = df3.to_pandas_on_spark()
psdf2.plot.bar(y = 'count(protocol_type)', x = 'protocol_type', color = 'protocol_type',
               text = 'count(protocol_type)')
```



```
In [15]: # Creating pivot table to check relationship between label and protocol type
df3 = df_sql.groupby('protocol_type').pivot('label_cat').agg(F.count('protocol_type')).orderBy('protocol_ty
psdf2 = df3.to_pandas_on_spark()
psdf2 = psdf2.fillna(0)
#psdf2.head()
df_pd = psdf2.set_index('protocol_type')
df_pd
```


	DOS	Probe	R2L	U2L	normal
protocol_type					
udp	979	1958	0	3	191348
tcp	1074241	26512	98	1077	768670
icmp	2808150	12632	0	0	12763

```
In [16]: # Plotting protocol type against our target variable
df_pd.plot.bar(barmode = 'stack')
```




```
In [17]: ### Plotting Bar chart for service type

df3 = df_sql.groupby('service').agg(F.count('service')).orderBy('service', ascending= True)
psdf2 = df3.to_pandas_on_spark()
psdf2.plot.bar(y = 'service', x = 'count(service)', color = 'service')
```




```
In [18]: # Creating pivot table to check relationship between label and service

df3 = df_sql.groupby('service').pivot('label_cat').agg(F.count('service')).orderBy('service', ascending= Fa
psdf2 = df3.to_pandas_on_spark()
psdf2 = psdf2.fillna(0)
#psdf2.head()
df_pd = psdf2.set_index('service')
df_pd.head(3)
```

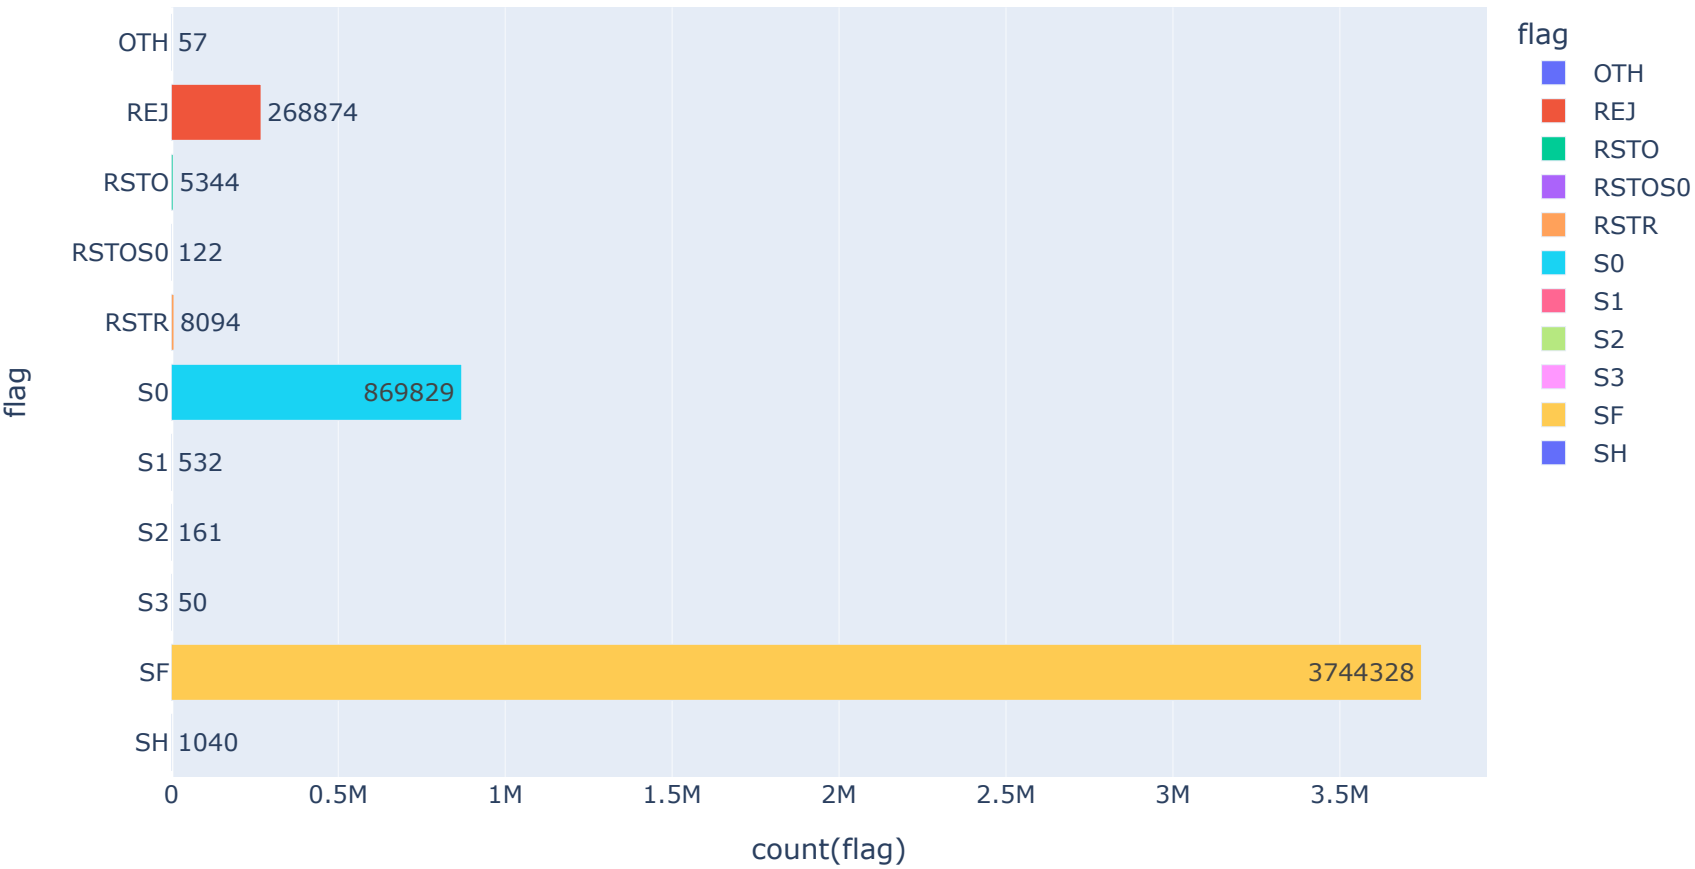
	DOS	Probe	R2L	U2L	normal
service					
whois	1042	31	0	0	0
vmnet	1041	12	0	0	0
uucp_path	1044	13	0	0	0

```
In [19]: # Plotting service against our target variable
df_pd.plot.bar(barmode = 'stack')
```



In [21]: *### Plotting Bar chart for flag type*

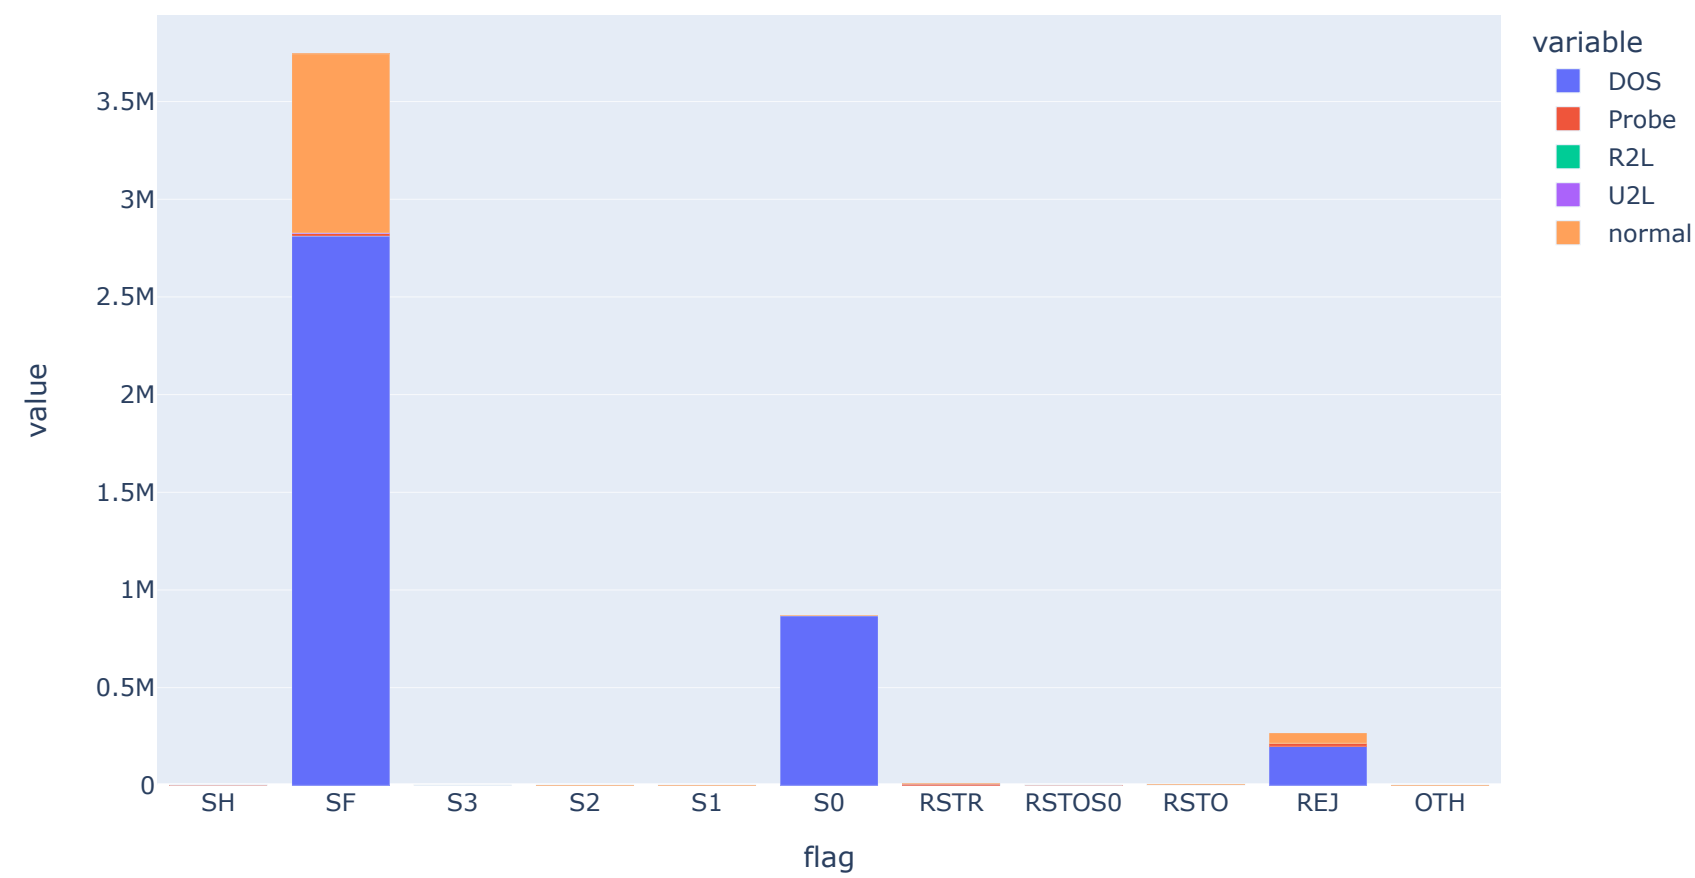
```
df3 = df_sql.groupby('flag').agg(F.count('flag')).orderBy('flag', ascending= True)
psdf2 = df3.to_pandas_on_spark()
psdf2.plot.bar(y = 'flag', x = 'count(flag)', color = 'flag', text = 'count(flag)')
```



```
In [22]: # Creating pivot table to check relationship between label and flag
df3 = df_sql.groupby('flag').pivot('label_cat').agg(F.count('flag')).orderBy('flag', ascending= False)
psdf2 = df3.to_pandas_on_spark()
psdf2 = psdf2.fillna(0)
#psdf2.head()
df_pd = psdf2.set_index('flag')
df_pd.head(3)
```

	DOS	Probe	R2L	U2L	normal
flag					
SH	0	1034	4	0	2
SF	2811235	14769	41	1075	917208
S3	0	1	2	1	46

```
In [23]: # Plotting flag against our target variable
df_pd.plot.bar(barmode = 'stack')
```



Identifying relationships between target & numerical independent variables by comparing means

```
In [25]: psdf4 = df_sql.to_pandas_on_spark()
psdf4.groupby('label_num').mean().T
```

	label_num	4	2	1	0	3
duration	585.213889	590.519464	0.000074	217.824724	25.091837	
src_bytes	283585.966667	109376.218870	707.446083	1477.846250	157.132653	
dst_bytes	966.800000	51350.530655	4.670547	3234.650111	822842.908163	
land	0.000000	0.000000	0.000005	0.000007	0.000000	
wrong_fragment	0.000000	0.000000	0.000818	0.000000	0.000000	
urgent	0.003704	0.000000	0.000000	0.000036	0.000000	
hot	7.673148	0.000462	0.001114	0.049535	1.091837	
num_failed_logins	0.000926	0.000097	0.000000	0.000099	0.571429	
logged_in	0.992593	0.002141	0.000568	0.719268	0.132653	
num_compromised	0.059259	0.000170	0.000548	0.038389	0.775510	
root_shell	0.024074	0.000000	0.000000	0.000310	0.061224	
su_attempted	0.000000	0.000000	0.000000	0.000184	0.010204	
num_root	0.039815	0.000170	0.000000	0.064970	1.112245	
num_file_creations	0.039815	0.000487	0.000000	0.005887	0.336735	
num_shells	0.006481	0.000000	0.000000	0.000363	0.040816	
num_access_files	0.003704	0.000000	0.000000	0.005131	0.071429	
num_outbound_cmds	0.000000	0.000000	0.000000	0.000000	0.000000	
is_host_login	0.000000	0.000000	0.000000	0.000002	0.000000	
is_guest_login	0.286111	0.000024	0.000000	0.003882	0.051020	
count	1.520370	171.770936	418.668688	8.159029	1.438776	
srv_count	1.331481	7.349496	369.633827	10.912790	12.887755	
serror_rate	0.004148	0.070904	0.223364	0.001483	0.105510	
srv_serror_rate	0.002444	0.073539	0.223360	0.001725	0.095408	
rerror_rate	0.002787	0.567238	0.052708	0.055941	0.500000	
srv_rerror_rate	0.002778	0.565939	0.052737	0.056206	0.507959	
same_srv_rate	0.993620	0.586935	0.743029	0.985257	1.000000	
diff_srv_rate	0.009241	0.407238	0.017760	0.018535	0.000000	
srv_diff_host_rate	0.011565	0.220565	0.000119	0.132494	0.046735	
dst_host_count	82.085185	169.054182	254.867245	148.498415	44.816327	
dst_host_srv_count	38.618519	52.952825	187.495878	202.014806	27.724490	
dst_host_same_srv_rate	0.738676	0.318464	0.735483	0.844879	0.931327	
dst_host_diff_srv_rate	0.021824	0.589753	0.018337	0.056500	0.001837	
dst_host_same_src_port_rate	0.658407	0.589580	0.722970	0.134940	0.368163	
dst_host_srv_diff_host_rate	0.098639	0.189274	0.000026	0.024340	0.010204	
dst_host_serror_rate	0.010352	0.070218	0.223382	0.002039	0.133571	
dst_host_srv_serror_rate	0.003537	0.073449	0.223338	0.001050	0.131429	

	label_num	4	2	1	0	3
dst_host_rerror_rate	0.005648	0.547672	0.052784	0.057784	0.477857	
dst_host_srv_rerror_rate	0.001407	0.564952	0.052707	0.056016	0.475510	

Observations:

- Duration of connection for attack is higher than normal except for DOS and R2L.
- Count of outbound commands in an ftp session are 0 for normal and attack connections.
- Wrong fragments in the connection is only present in DOS attack.

Correlation Matrix

```
In [28]: # Creating Correlation Matrix
psdf_corr = df_sql.to_pandas_on_spark()
corrmm=psdf_corr[numerical_variables].corr()
corrmm
```

	duration	src_bytes	dst_bytes	land	wrong_fragment	urgent	hot	num_failed_log
duration	1.000000	4.122055e-02	0.020392	-0.000160	-0.001012	3.765465e-03	0.004450	0.007412
src_bytes	0.041221	1.000000e+00	0.000239	-0.000005	-0.000027	-8.677487e-08	0.000782	-0.000007
dst_bytes	0.020392	2.393376e-04	1.000000	-0.000004	-0.000026	1.645208e-04	0.000126	0.000632
land	-0.000160	-4.659182e-06	-0.000004	1.000000	-0.000036	-2.638263e-06	-0.000063	-0.000010
wrong_fragment	-0.001012	-2.714937e-05	-0.000026	-0.000036	1.000000	-1.670585e-05	-0.000402	-0.000066
urgent	0.003765	-8.677487e-08	0.000165	-0.000003	-0.000017	1.000000e+00	0.003591	0.031005
hot	0.004450	7.822053e-04	0.000126	-0.000063	-0.000402	3.590666e-03	1.000000	0.004475
num_failed_logins	0.007412	-6.921906e-06	0.000632	-0.000010	-0.000066	3.100532e-02	0.004475	1.000000
logged_in	-0.020624	1.998519e-04	0.002119	-0.000979	-0.006197	2.534183e-03	0.064579	0.001792
num_compromised	0.027126	4.832481e-06	0.001307	-0.000005	-0.000032	1.767953e-02	0.002688	0.019542
root_shell	0.026378	-5.620801e-06	0.000988	-0.000020	-0.000125	8.908420e-02	0.017916	0.023673
su_attempted	0.052088	-4.122588e-06	0.001204	-0.000011	-0.000069	1.330230e-01	0.001926	0.069186

Creating a Pipeline to encode the qualitative variables which are of String datatype namely protocol_type, service, flag and label into variables of datatype Double by using StringIndexer Transformers.

This step is important as Logistic Regression and NaiveBayes classifiers needs to be trained on numerical values.

<https://spark.apache.org/docs/latest/api/python/reference/api/pyspark.ml.feature.StringIndexer.html>
(<https://spark.apache.org/docs/latest/api/python/reference/api/pyspark.ml.feature.StringIndexer.html>)

```
In [26]: # Defining Categorical Variables
         qualitative_variables = ["protocol_type", "service", "flag"]

         # Using String Indexer to encode Categorical variables to Numerical Variables
         indexers = [StringIndexer(inputCol=column, outputCol=column + "_num") for column in qualitative_variables]

         #Creating a pipeline
         pipeline = Pipeline(stages=indexers)
         #Transforming our spark dataframe
         df_sql = pipeline.fit(df_sql).transform(df_sql)

         # Excluding Non-Numeric Attributes
         not_needed = qualitative_variables + ["label", "label_num", "label_cat"]
         #print(not_needed)
```

```
In [27]: numerical_variables = [col for col in df_sql.columns if col not in not_needed]
print(numerical_variables)
df_sql.printSchema()
```

```
['duration', 'src_bytes', 'dst_bytes', 'land', 'wrong_fragment', 'urgent', 'hot', 'num_failed_logins', 'logged_in', 'num_compromised',
'root_shell', 'su_attempted', 'num_root', 'num_file_creations', 'num_shells', 'num_access_files', 'num_outbound_cmds', 'is_host_login',
'is_guest_login', 'count', 'srv_count', 'serror_rate', 'srv_serror_rate', 'rerror_rate', 'srv_rerror_rate', 'same_srv_rate', 'diff_srv_r
ate', 'srv_diff_host_rate', 'dst_host_count', 'dst_host_srv_count', 'dst_host_same_srv_rate', 'dst_host_diff_srv_rate', 'dst_host_same_s
rc_port_rate', 'dst_host_srv_diff_host_rate', 'dst_host_serror_rate', 'dst_host_srv_serror_rate', 'dst_host_rerror_rate', 'dst_host_srv_
rerror_rate', 'protocol_type_num', 'service_num', 'flag_num']
root
|-- duration: integer (nullable = true)
|-- protocol_type: string (nullable = true)
|-- service: string (nullable = true)
|-- flag: string (nullable = true)
|-- src_bytes: integer (nullable = true)
|-- dst_bytes: integer (nullable = true)
|-- land: integer (nullable = true)
|-- wrong_fragment: integer (nullable = true)
|-- urgent: integer (nullable = true)
|-- hot: integer (nullable = true)
|-- num_failed_logins: integer (nullable = true)
|-- logged_in: integer (nullable = true)
|-- num_compromised: integer (nullable = true)
|-- root_shell: integer (nullable = true)
|-- su_attempted: integer (nullable = true)
|-- num_root: integer (nullable = true)
|-- num_file_creations: integer (nullable = true)
|-- num_shells: integer (nullable = true)
|-- num_access_files: integer (nullable = true)
|-- num_outbound_cmds: integer (nullable = true)
|-- is_host_login: integer (nullable = true)
|-- is_guest_login: integer (nullable = true)
|-- count: integer (nullable = true)
|-- srv_count: integer (nullable = true)
|-- serror_rate: double (nullable = true)
|-- srv_serror_rate: double (nullable = true)
|-- rerror_rate: double (nullable = true)
|-- srv_rerror_rate: double (nullable = true)
|-- same_srv_rate: double (nullable = true)
|-- diff_srv_rate: double (nullable = true)
|-- srv_diff_host_rate: double (nullable = true)
|-- dst_host_count: integer (nullable = true)
|-- dst_host_srv_count: integer (nullable = true)
|-- dst_host_same_srv_rate: double (nullable = true)
|-- dst_host_diff_srv_rate: double (nullable = true)
|-- dst_host_same_src_port_rate: double (nullable = true)
|-- dst_host_srv_diff_host_rate: double (nullable = true)
|-- dst_host_serror_rate: double (nullable = true)
|-- dst_host_srv_serror_rate: double (nullable = true)
|-- dst_host_rerror_rate: double (nullable = true)
|-- dst_host_srv_rerror_rate: double (nullable = true)
|-- label: string (nullable = true)
|-- label_num: integer (nullable = false)
```

```
|-- label_cat: string (nullable = false)
|-- protocol_type_num: double (nullable = false)
|-- service_num: double (nullable = false)
|-- flag_num: double (nullable = false)
```

Now we will use Spark's Vector Assembler to create a feature vector for each numerical variable as we need to combine all the input columns into a single vector which would essentially act as the input feature for the classifiers.

```
In [29]: df_assembler = VectorAssembler(inputCols=numerical_variables, outputCol="features")
df_sql = df_assembler.transform(df_sql)
#df_sql.printSchema()
```

We can see a features column has been created in the dataframe schema

Now we will create a new dataframe *ml_df* containing two columns - features and label_num. Furthermore, we will randomly split it into the training dataset and testing dataset in 75:25 ratio respectively.

<https://medium.com/udemy-engineering/pyspark-under-the-hood-randomsplit-and-sample-inconsistencies-examined-7c6ec62644bc> (<https://medium.com/udemy-engineering/pyspark-under-the-hood-randomsplit-and-sample-inconsistencies-examined-7c6ec62644bc>) - The signature function of randomSplit() includes a weight list and a seed specification. The weight list is to specify the number of splits and percentage (approximate) in each and the seed is for reproducibility. The ratio is approximate due to the nature of how it is calculated.

```
In [30]: # Selecting vectorized features column and label_num
ml_df = df_sql.select(["features", "label_num"])
ml_df.printSchema()

train_set, test_set = ml_df.randomSplit([0.75, 0.25], seed=3000)
print("Training dataset count: " + str(train_set.count()))
print("Test dataset count: " + str(test_set.count()))

root
 |-- features: vector (nullable = true)
 |-- label_num: integer (nullable = false)

Training dataset count: 3675127
Test dataset count: 1223304
```

```
In [31]: # Selecting vectorized features column and label_num
ml_df = df_sql.select(["features","label_num"])
ml_df.printSchema()

train_set, test_set = ml_df.randomSplit([0.75, 0.25], seed=2000)
print("Training dataset count: " + str(train_set.count()))
print("Test dataset count: " + str(test_set.count()))

root
|-- features: vector (nullable = true)
|-- label_num: integer (nullable = false)

Training dataset count: 3672779
Test dataset count: 1225652
```

From above two blocks of code we can verify the random split by checking the dataset shape

Building Machine Learning Models

KPI(Key Performance Indicator)

- Accuracy
- Weighted Precision
- Weighted Recall
- F1 Score

```
In [35]: metrics = ["accuracy", "weightedPrecision", "weightedRecall", "f1"]
```

Evaluation

Logistic Regression

```
In [88]: # Logistic Regression model
lr = LogisticRegression(maxIter=20, regParam=0.3, elasticNetParam=0.8, featuresCol="features",
                        labelCol="label_num", family="multinomial")

lrlist = []

print("\nLogistic Regression Model Evaluation:")
print("-"*30)
lr_model = lr.fit(train_set)

# make predictions
predictions = lr_model.transform(test_set)
predictions.cache()

evaluator = MulticlassClassificationEvaluator(labelCol="label_num", predictionCol="prediction")

for m in metrics:
    evaluator.setMetricName(m)
    metric = evaluator.evaluate(predictions)
    print("{name} = {value:.2f}".format(name=m, value=metric))
    lrlist.append(metric)

predictions.show(5)
```

```
Logistic Regression Model Evaluation:
-----
accuracy = 0.79
weightedPrecision = 0.63
weightedRecall = 0.79
f1 = 0.70
+-----+-----+-----+-----+-----+
|          features|label_num|    rawPrediction|    probability|prediction|
+-----+-----+-----+-----+-----+
|(41,[0,1,2,6,8,10...|      0|[4.04992757798011...|[0.38493120112223...|    1.0|
|(41,[0,1,2,6,8,18...|      0|[4.04952262397637...|[0.38452734500661...|    1.0|
|(41,[0,1,2,6,8,18...|      0|[4.04972510097824...|[0.38472925408982...|    1.0|
|(41,[0,1,2,6,8,18...|      0|[4.04972510097824...|[0.38472925408982...|    1.0|
|(41,[0,1,2,6,8,18...|      0|[4.04972510097824...|[0.38472925408982...|    1.0|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

Decision Tree

```
In [89]: # Decision Tree model

dt = DecisionTreeClassifier(labelCol="label_num", featuresCol="features",  maxBins=70)

dtlist = []

print("\nDecision Tree Model Evaluation:")
print("{:-<30}".format(""))
dt_model = dt.fit(train_set)

# make predictions
predictions = dt_model.transform(test_set)
predictions.cache()

evaluator = MulticlassClassificationEvaluator(labelCol="label_num", predictionCol="prediction")

for m in metrics:
    evaluator.setMetricName(m)
    metric = evaluator.evaluate(predictions)
    print("{name} = {value:.2f}".format(name=m, value=metric))
    dtlist.append(metric)

predictions.show(5)
```

```
Decision Tree Model Evaluation:
-----
accuracy = 1.00
weightedPrecision = 1.00
weightedRecall = 1.00
f1 = 1.00
+-----+-----+-----+-----+
|          features|label_num|    rawPrediction|    probability|prediction|
+-----+-----+-----+-----+
|(41,[0,1,2,6,8,10...|      0|[708161.0,706.0,4...|[0.99177212353438...|      0.0|
|(41,[0,1,2,6,8,18...|      0|[708161.0,706.0,4...|[0.99177212353438...|      0.0|
|(41,[0,1,2,6,8,18...|      0|[708161.0,706.0,4...|[0.99177212353438...|      0.0|
|(41,[0,1,2,6,8,18...|      0|[708161.0,706.0,4...|[0.99177212353438...|      0.0|
|(41,[0,1,2,6,8,18...|      0|[708161.0,706.0,4...|[0.99177212353438...|      0.0|
+-----+-----+-----+-----+
only showing top 5 rows
```

Random Forest Classifiier

```
In [90]: # Random Forest model

rf = RandomForestClassifier(labelCol="label_num", featuresCol="features", numTrees=20, maxBins=70)

rflist = []

print("\nRandom Forest Classifier Model Evaluation:")
print("{:-<30}".format(""))
rf_model = rf.fit(train_set)

# make predictions
predictions = rf_model.transform(test_set)
predictions.cache()

evaluator = MulticlassClassificationEvaluator(labelCol="label_num", predictionCol="prediction")

for m in metrics:
    evaluator.setMetricName(m)
    metric = evaluator.evaluate(predictions)
    print("{name} = {value:.2f}".format(name=m, value=metric))
    rflist.append(metric)

predictions.show(5)
```

```
Random Forest Classifier Model Evaluation:
-----
accuracy = 1.00
weightedPrecision = 1.00
weightedRecall = 1.00
f1 = 1.00
+-----+-----+-----+-----+-----+
|          features|label_num|    rawPrediction|    probability|prediction|
+-----+-----+-----+-----+-----+
|(41,[0,1,2,6,8,10...|      0|[16.9687825289452...|[0.84843912644726...|      0.0|
|(41,[0,1,2,6,8,18...|      0|[19.6540507623080...|[0.98270253811540...|      0.0|
|(41,[0,1,2,6,8,18...|      0|[19.5035822798782...|[0.97517911399391...|      0.0|
|(41,[0,1,2,6,8,18...|      0|[19.6420975918839...|[0.98210487959419...|      0.0|
|(41,[0,1,2,6,8,18...|      0|[19.6520642213605...|[0.98260321106802...|      0.0|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

Naive Bayes Multinomial

```
In [91]: # Naive Bayes Multinomial
nb = NaiveBayes(labelCol="label_num", featuresCol="features", smoothing=1.0, modelType="multinomial")

nblist = []

print("\nNaive Bayes Multinomial Model Evaluation:")
print("{:-<30}".format(""))
nb_model = nb.fit(train_set)

# make predictions
predictions = nb_model.transform(test_set)
predictions.cache()

evaluator = MulticlassClassificationEvaluator(labelCol="label_num", predictionCol="prediction")

for m in metrics:
    evaluator.setMetricName(m)
    metric = evaluator.evaluate(predictions)
    print("{name} = {value:.2f}".format(name=m, value=metric))
    nblist.append(metric)

predictions.show(5)
```

```
Naive Bayes Multinomial Model Evaluation:
-----
accuracy = 0.93
weightedPrecision = 0.95
weightedRecall = 0.93
f1 = 0.93
+-----+-----+-----+-----+-----+
|          features|label_num|    rawPrediction|    probability|prediction|
+-----+-----+-----+-----+-----+
|(41,[0,1,2,6,8,10...|      0|[-384791.27476208...|[0.0,0.0,0.0,1.0,...|      3.0|
|(41,[0,1,2,6,8,18...|      0|[-1136.0742561634...|[1.0,0.0,2.097648...|      0.0|
|(41,[0,1,2,6,8,18...|      0|[-4548.3207874904...|[1.0,0.0,0.0,0.0,...|      0.0|
|(41,[0,1,2,6,8,18...|      0|[-1706.7583981789...|[1.0,0.0,0.0,0.0,...|      0.0|
|(41,[0,1,2,6,8,18...|      0|[-1325.6372574940...|[1.0,0.0,3.609994...|      0.0|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

Comparing our models


```
In [92]: l = [round(item, 2) for item in lrlist]
l.insert(0, 'Logistic Regression')
d = [round(item, 2) for item in dtlist]
d.insert(0, 'Decision Tree')
r = [round(item, 2) for item in rflist]
r.insert(0, 'Random Forest Classifier')
n = [round(item, 2) for item in nblist]
n.insert(0, 'Naive Bayes')

data = [l,d,r,n]

#giving column names of dataframe
columns = ["Models","Accuracy", "WeightedPrecision", "WeightedRecall","F1"]
#creating a dataframe
dataframe = spark.createDataFrame(data, columns)

#show data frame
dataframe.show()
```

	Models	Accuracy	WeightedPrecision	WeightedRecall	F1
	Logistic Regression	0.79	0.63	0.79	0.7
	Decision Tree	1.0	1.0	1.0	1.0
	Random Forest Cla...	1.0	1.0	1.0	1.0
	Naive Bayes	0.93	0.95	0.93	0.93

```
In [93]: psdf_results = dataframe.to_pandas_on_spark()
psdf_results = psdf_results.set_index('Models')
psdf_results.head()
```

	Accuracy	WeightedPrecision	WeightedRecall	F1
Models				
Logistic Regression	0.79	0.63	0.79	0.70
Decision Tree	1.00	1.00	1.00	1.00
Random Forest Classifier	1.00	1.00	1.00	1.00
Naive Bayes	0.93	0.95	0.93	0.93

```
In [94]: | psdf_results.plot.bar(barmode='group').update_xaxes(categoryorder="total ascending")
```

Saving our model file

Saving Decision Tree Model file using pickle as it has the highest accuracy.

<https://suziepyspark.blogspot.com/2021/03/using-pickle-to-save-model-in-incorta.html>
(<https://suziepyspark.blogspot.com/2021/03/using-pickle-to-save-model-in-incorta.html>)

```
In [380]: #!/pip install prophet
          from pyspark.sql.functions import *
          from pyspark.sql.types import *
          from prophet import Prophet
          import pickle

          import pickle
          pkl_path = "model.pkl"
          with open(pkl_path, "wb") as f:
              #Saving our decision tree model
              pickle.dump(dt_model, f)

          # save the model
          print("*** Data Saved ***")
```

*** Data Saved ***

```
In [384]: import pickle
          model=pickle.load(open('model.pkl', 'rb'))

          # Parameters for DOS
          model.predict([[1,0,0,1000,0.06,0.00,0.04,10,0,0,21,0,0.04,0.00,0]])

          array([1.])
```

```
In [382]: # Parameters for probe
          model.predict([[0,0,0,0,0.06,0.00,0.04,10,0,0,21,0,0.04,0.00,0]])

          array([2.])
```

References

- <https://sparkbyexamples.com/pyspark/different-ways-to-create-dataframe-in-pyspark/>
- <https://github.com/biagiom/spark-network-traffic-classifier>
- <https://unsplash.com/>
- <https://www.barracuda.com/glossary/intrusion-detection-system>
- <https://www.ecb.torontomu.ca/~bagheri/papers/cisda.pdf>

